

Library Supply and Search Satisfaction

Seeding, creator acquisition, reader demand, and platform support

Sponsor	Adam Kershner, Co-Founder
Project Manager	Nithila
Prepared	23 July 2026
Version	v4.0 (Revised on sponsor feedback, for team review)
Informed by	Sponsor feedback 23 July, Women Creators New Market Research, Scaling Kahana charter v2.0, about.kahana.io

1 Vision

Before the objectives, the feeling we are building for.

You are on campus and a friend texts asking if you want to pregame for the party. You went out last night, and the weekend before. You say thanks, but you are headed to the library.

Inside, it is quiet and warm. People are curled into chairs in comfortable hoodies, deep in their work, laptops and books open. Shelves rise around you holding more than you could read in a lifetime. The size of it makes you feel small in a good way. Surrounded by centuries of collected knowledge and other focused people, you become curious and open. You fall into the reading. Hours pass without your noticing. When you step outside into clear night air, ideas are connecting in your head and you feel better than when you walked in.

Instagram, TikTok, YouTube and Netflix are the virtual party. They are always there when you want a dopamine hit, and there is nothing wrong with a party. But eventually you graduate and lose your library card, or you move somewhere without a good public library. Kahana is the digital library with Aura: a place you can always go to fall into focus, feel productive, and be surrounded by other learners. We are building the room people deserve for thinking, learning, and becoming better.

Everything in this charter serves that experience. A reader who searches and finds nothing does not get that feeling. That is the problem we are solving.

2 Purpose

Kahana is The Digital Library with Aura. Readers explore hubs of digital artifacts, mainly eBooks and videos, and spend a daily budget of Aura to recognize work worth noticing. The library only delivers the experience in section 1 when a search returns something worth opening.

As of 21 July 2026, the platform holds 1,272 listed authors, 1,550 hubs, and roughly 17.9k files across 39 searchable topics. Coverage is uneven, so a meaningful share of searches returns thin or empty results. That is the constraint this project exists to remove.

Four workstreams run in parallel. Seeding fills gaps directly with openly licensed material. Creator acquisition recruits people to publish original work, eBooks first and video second. Reader demand research finds out what people actually search for and fail to find. Platform and technical support ships product changes that make all three faster. The project is judged on whether searches succeed, not on how many items we added.

3 Success Criteria and KPIs

The sponsor has set the standard for this project: we are solving for reader outcomes, not supply counts. A library can add ten thousand files and still fail if people search and leave empty-handed.

Product heuristic

For our top personas, at least 80 percent of their typical searches should yield multiple relevant results across files, hubs and authors, before we expect strong organic retention and word of mouth.

Primary success criteria

- CURR, current user retention rate, the chance a user returns this week having been active the past two weeks.
- DAUs and cohort retention on Explore and hub access.
- Customer delight: session depth, return visits, Aura given, and qualitative feedback.
- Reduction in search sessions returning zero or thin results.

New KPIs this project owns

- Search success rate: the percentage of searches that lead to a meaningful engagement, meaning a view, save, share, Aura grant, or the start of a borrow or purchase. This is the headline number for the project.
- Liquidity by niche: for priority segments, how often a searcher finds enough variety and quality to feel the library has what they need. Measured per topic cluster rather than platform-wide, since an average hides the gaps.
- Retention and referrals among high-search-satisfaction users: whether people who find what they want return more often and recommend the platform. This is the earliest signal of organic growth and network effects.

Secondary, expected to follow from the above

- MRR and ARR from Growth subscriptions and hub revenue.
- GMV from onboarded creators at the 5 percent take rate.

Operational tracking, not success criteria

Counts of seeded titles, creators onboarded, hubs published, and Aura on new hubs are still tracked, but as workload and pacing indicators for the pods rather than as measures of success. They tell us whether the machine is running. They do not tell us whether it is working. The dashboard shows them in a separate section from the KPIs above, so the distinction stays visible.

Instrumentation note views, saves, purchases, Aura grants and shares are all measurable through Mixpanel and Firebase. Search success rate requires linking a search event to the engagement that follows it, which is an instrumentation task for the technical team and a dependency for reporting.

4 Reader Demand: How We Know What to Fill

Topic rankings based on AI research are a starting hypothesis, not evidence. If we seed purely to the ranked list, we will fill the library with what we assume women want rather than what they actually search for and fail to find.

Search panel research

- Recruit women readers for virtual or in-person sessions and ask them to share their screen while using the library to find something they would genuinely save, read, or watch.
- Log every query that fails to return actionable results. The failed searches are the deliverable, not the successful ones.
- Feed those gaps directly into seeding and outreach priorities. If a reader wants a Filipino recipe book and the library has nothing, that becomes a sourcing target.
- Re-run the panel periodically so priorities follow real demand as coverage changes.

Standing search set

The following ten searches, drafted by the sponsor in the way real users phrase them, form a repeatable test. Run them weekly, score whether each return multiple relevant results, and track the trend. This gives a fast proxy for search success while full instrumentation is built.

#	Search
1	Empowering fiction about women in business and leadership
2	Psychology eBooks on relationships, attachment styles, and emotional health
3	Biographies of inspiring women in history, politics, and science
4	Modern literature and fiction about motherhood, family, and identity
5	Guides on women's health and medicine, including mental health and self-care
6	Books on society and culture focusing on gender equality and social justice
7	Digital marketing and management eBooks for women entrepreneurs
8	Art and design books featuring women creators and creative careers
9	Travel memoirs and adventure stories written by women
10	Religion and philosophy books exploring spirituality, purpose, and meaning from a women's perspective

5 Topics

Kahana now supports 39 topics that authors can apply to their hubs, so they appear in search results. This replaces the earlier 16 and is the authoritative taxonomy for coverage. The ordering below reflects the sponsor's ranking by likely appeal to a broad adult women audience, and is used to sequence effort, not to exclude anything.

Tier	Topics	Use
1 to 10	Fiction, Literature, Biography, History, Psychology, Society, Culture, Travel, Art, Film	Lead here. Seeding and outreach concentrate on these first, adjusted by what the search panel shows is missing.
11 to 20	Music, Crafts, Design, Games, Recreation, Drama, Poetry, Education, Business, Marketing	Second wave. Begin once tier one has depth, or earlier if panel research shows demand.
21 to 30	Management, Medicine, Biology, Science, Technology, Computing, Information, Knowledge, Language, Mythology	Fill opportunistically, especially where a creator or open-access source is already available.
31 to 39	Religion, Philosophy, Politics, Economics, Law, Geography, Engineering, Linguistics, Sports	Lowest priority for deliberate effort. Accept organic and creator-driven contributions.

Ranking is a hypothesis about appeal. Where the search panel contradicts it, the panel wins.

6 Scope

In scope

- Sourcing, license-checking, and publishing openly licensed titles from team accounts to fill gaps.
- Sourcing, qualifying, and onboarding new creators, eBooks first and video second.
- Reactivating existing creators who published once and stopped.
- Reader demand research through search panels and translating failed searches into sourcing priorities.
- Changes to the platform itself. The technical team can ship updates multiple times per week and is available for the duration. This includes search and discovery improvements, onboarding and upload flow, tagging, previews, import tooling, instrumentation, and changes to how Aura behaves where they serve the objectives.
- Instrumentation in Mixpanel and Firebase to measure search success, engagement per search, and retention among satisfied searchers.
- Outreach infrastructure: Google Workspace accounts across the Kahana domains, sender warm-up, and sequenced outreach from the founder.
- The partnership offer: white glove support, Growth plan discount, featured placement, and Aura-driven exposure.

Out of scope

- Licensing catalogs from publishing houses. A B2B motion needing legal review and contract terms, more viable once we can show reader demand.
- Reader-side paid acquisition campaigns. This project fixes what readers find, not how many arrive.

- Content formats beyond eBooks, documents and video in this run.

Nothing in the product is off limits. Where a platform change would make searches succeed more often, it is in scope and goes to the technical backlog.

7 Workstream A: Seeding

Purpose is coverage in the topics where readers search and find nothing.

- Use openly licensed material only. Sources include the Library of Congress Open Access Books collection, the Directory of Open Access Books, and Project Gutenberg.
- Verify license per title rather than trusting a collection label. Gutenberg in particular mixes public domain works with titles held under specific permissions, so titles need individual checking.
- Do not mass-download. Bulk scraping risks IP blocking and, in some cases, legal exposure. Acquire deliberately and at a human pace.
- Prefer permissive licenses. Treat Noncommercial terms with caution given the platform monetizes elsewhere.
- Honour attribution on every title.
- Publish from clearly identified team accounts, labelled as library collection, so readers can distinguish curated public material from creator work.
- Legal due diligence is owned by the sponsor and must clear before seeding at volume.

Seeding is bounded and targeted. Volume is set by what the search panel and topic gaps require, not maximized. Seeded hubs are reported separately from creator work in every update.

8 Workstream B: Creator Acquisition

Phase one targets eBook authors, phase two video creators, following the sponsor's view that adding a book to a library is an intuitive ask while video needs more explanation of what belongs.

The workstream runs in two stages. Sourcing builds a reliable, rights-qualified list. Outreach works that list. The two overlap but sourcing leads, because outreach into a thin or unqualified list wastes both the team's time and the sending reputation of new domains.

Stage 1: Sourcing and list building

This is deliberate research work, not a side task before the real job. It runs from week one and continues throughout, staying ahead of outreach capacity.

- Work the channels per segment in the table below, building a prospect record for each name: source channel, segment, platform links, audience size, what they currently sell, and contact route.
- Apply the rights qualification in the next section before a name enters the outreach queue. An unqualified name is a lead, not a prospect.
- Verify contact details. A list with stale or guessed email addresses damages deliverability as much as bad targeting.
- Use the technical team for automation where it helps: scraping seller directories, enriching records, and deduplicating. Raise it with Divya rather than doing it by hand at volume.
- Track list health weekly: total qualified prospects, new added, and how many weeks of outreach capacity the list covers.

Checkpoint before scaling outreach at volume

Outreach moves beyond small test batches only when the list holds at least 200 rights-qualified prospects with verified contact details, and the sending domains have completed warm-up. Small test sends of 10 to 20 can start earlier to sharpen the message and should.

Sourcing and domain warm-up both take weeks, so they run in parallel from week one. By the time senders are trusted enough for volume, the list should be ready to feed them.

Stage 2: Outreach and onboarding

- Sequenced sends from warmed senders, personalized by segment, with the offer of white glove support, a Growth plan discount, and featured placement.
- Track every prospect through contacted, replied, call scheduled, onboarded, with the decline reason captured on every no.
- Onboard using the playbook: profile, first hub, import of existing catalog, Stripe Connect for paid work, eBook tagging, and the daily Aura habit.
- Time featured placement to launch and prompt the creator's own audience to give Aura, so new work does not publish into silence.

Rights qualification

Tier	Who	Where to find them
Priority 1	Self-published authors who are not Amazon-exclusive. They control their digital rights and can accept immediately.	Etsy, Gumroad, Ko-fi, Payhip, Substack, and personal sites. Selling PDFs outside Amazon is itself evidence they are not exclusive.
Priority 2, parallel test	Authors who control something even if their main title is unavailable: reverted backlist rights, self-published side projects, companion workbooks.	Same channels, plus Goodreads and author networks. Requires one extra qualifying question.
Excluded, phase one	Authors in Amazon KDP Select, who have agreed to digital exclusivity, and traditionally published authors whose publisher holds eBook rights.	Practical filter: a title in Kindle Unlimited means the author is in KDP Select and therefore exclusive.

Outreach infrastructure

Outreach runs from the founder personally, across Google Workspace accounts on the Kahana domains (kahana.app, kahana.club, kahana.io, kahana.llc, kahana.vip, kahana.work). Standard practice is to start with five sender addresses and send five to ten emails per day per sender, building domain authority and deliverability before scaling volume. The offer is white glove support, a Growth plan discount, and featured placement.

Sender warm-up takes weeks, so set the accounts up in the first fortnight even though outreach volume stays low initially. Deliverability is the constraint, not list size.

Segments

Sourcing segments and channels come from the Women Creators New Market Research. Women-first is a sourcing guideline for building the prospect list, not a reported metric, since the platform does not capture gender at profile creation.

Segment	Priority	Where to source
Writing & Publishing	High	Substack, BookTok and Wattpad communities, non-exclusive KDP authors, Alliance of Independent Authors, NaNoWriMo alumni
Education & Coaching	High	Teachable, Thinkific and Kajabi directories, LinkedIn coaching communities, alumni and EdTech networks
Arts, Illustration & Design	High	Creative Market, Etsy and Gumroad seller lists, Behance and Dribbble, Procreate and Canva groups
Parenting, Family & Home	Medium-High	Parenting forums and blogs, Etsy printable shops, Facebook groups and Pinterest, homeschool co-ops
Productivity & Planners	Medium-High	Notion template marketplaces, productivity YouTube and TikTok, Etsy planner top sellers
Spirituality & Mindfulness	Medium-High	Wellness platform creator lists, astrology and tarot communities, journaling Etsy sellers
Finance & Business	Medium	Personal finance blogs and podcasts, women-in-business associations, LinkedIn finance groups
Photography & Presets	Medium	Photography creator communities in lifestyle, portrait and wedding niches

9 Workstream C: Platform and Technical

The technical team supports the whole project and can ship multiple times per week. Divya Emmadi owns this workstream end to end: intake, prioritisation, backlog in Linear, and weekly triage with the technical team. Anyone in the project raises a request directly to Divya rather than routing it through the Project Manager.

Anticipated technical work

- Instrumentation linking search events to subsequent engagement, so search success rate can be reported. This is the highest priority technical item because the project's headline KPI depends on it.
- Search and discovery improvements where panel research shows queries failing for reasons other than missing content, for example poor matching or unhelpful ranking.
- File-level eBook tagging, and any bulk or assisted tagging that helps existing content surface.
- Upload, import and onboarding improvements that reduce time to first published hub.
- Preview or sampling on paid hubs, given observed view-to-purchase conversion is low.
- Automation support for seeding: metadata extraction, license capture, and bulk hub creation from approved sources.
- Changes to Aura behavior where they help new work gain early visibility.

Requests should arrive with the reader problem they solve and the KPI they move, so the technical team can sequence by impact rather than by who asked.

10 Deliverables

- Search panel research programme: recruitment, session protocol, and a running failed-query log that feeds sourcing priorities directly.
- Search success instrumentation in Mixpanel and Firebase, with a dashboard separating KPIs from operational counts.
- Weekly standing search set results across the ten sponsor queries, scored and trended.
- Topic gap audit across the 39 topics, refreshed as coverage changes.
- Seeding playbook: approved sources, per-title license procedure, attribution and labelling standard, and legal clearance record.
- Creator sourcing plan by segment and channel, with an agreed prospect record format and a weekly list health report.
- Creator prospect sheet and rights-qualification rubric. Gate for volume outreach is 200 rights-qualified prospects with verified contacts, building to at least 500 by end of week 8. Fields include source channel, segment, rights tier, contact status, and decline reason.
- Outreach infrastructure: Workspace accounts across the Kahana domains, warmed senders, and sequenced templates.
- Onboarding playbook covering profile, first hub, import, Stripe Connect, tagging, and Aura habit.
- Technical backlog in Linear, with each request stating the reader problem and target KPI.

11 Milestones and Timeline

Phase	Milestone	Target window
Setup	Charter reviewed by the full team, pods confirmed, topic gap audit complete, legal clearance route agreed, Workspace domains provisioned and warming started	Weeks 1 to 2
Instrument	Search success tracking live in Mixpanel and Firebase, dashboard split between KPIs and operational counts, baseline search success rate established	Weeks 2 to 5
Research	First search panel sessions run, failed-query log producing sourcing priorities	Weeks 3 to 6
Source	Prospect list built and rights-qualified to the 200 gate, senders warming, small test sends running	Weeks 1 to 6
Fill and recruit	Seeding underway in tier one topics, outreach at volume from warmed senders, first creators onboarded	Weeks 6 to 12
Scale	Outreach at steady cadence, reactivation campaign running, technical improvements shipping weekly against panel findings	Weeks 10 to 18

Video phase	Video creator guidance published, first video outreach batch	Weeks 14 to 20
Optimize	Priorities re-cut from panel and search data, gaps closed where success rate lags	Weeks 16 to 22
Review	Search success rate and retention measured against the 80 percent heuristic, decision to scale, adjust, or stop	Weeks 23 to 24

12 Team and Responsibilities

Twenty contributors across delivery and technical groups. Each member is expected to allocate at least 20 hours per week, Monday to Friday, barring illness or emergency.

Delivery group

Pod	Members	Responsibility
Pod A, Seeding	Tanmay Dave (lead), Hasitha Sigatapu, Ravi Chandra Dasari, Rohan Varshney, Khushmeet Arora	Source and license-check openly licensed titles, publish and quality-check seeded hubs, fill topic gaps identified by research.
Pod B, Creator outreach	Nithish Sampath (lead), Hasan Bohra, Sudhire Karunakaran, Sriram Reddy, Rashmila Mitra, Rashmi Kadwani	Stage 1 sourcing and list building, stage 2 outreach and onboarding, reactivation of existing creators. Rashmi also covers design across both pods.
Shared, Research and analytics	Niraj Patil, Jaideep Kulkarni	Search panel sessions and failed-query log, baseline and KPI reporting, dashboard, topic gap audit.

Technical group

Available to the whole project, shipping multiple times per week from the Linear backlog.

Member	Focus
Divya Emmadi	Product Manager. Owns the technical workstream: request intake, backlog prioritisation, weekly triage, and delivery.
Gayatri Talluri	Engineering and analytics.
Shreyans Gupta	Analytics.
Veda Gundu	Analytics and engineering. India time zone, so schedule handovers accordingly.
Archit Gupta	Engineering.
Ashwin John	Engineering and product.
Arpita Dev	Product and engineering.

How the work is run

Each pod and workstream is self-managing. Leads own their own planning, task assignment, and weekly status. The Project Manager does not assign individual tasks or chase updates. Status flows up on a fixed weekly rhythm, and anything blocked escalates to the Project Manager, who takes it to the Sponsor if it needs a decision above the team.

Owner	What they own outright
Tanmay Dave, Pod A lead	Seeding plan and pacing, task assignment within Pod A, weekly Pod A status.
Nithish Sampath, Pod B lead	Outreach plan and pacing, prospect and onboarding pipeline in Linear, task assignment within Pod B, weekly Pod B status.
Divya Emmadi, Technical	Technical backlog end to end: intake, prioritisation, triage, and delivery against the project's KPIs.
Niraj Patil and Jaideep Kulkarni, Analytics	Panel research and the failed-query log, instrumentation spec, and the live KPI dashboard that serves as the project's status.
Rashmi Kadwani, Pod B	Outreach copy and creative within Pod B, plus design across both pods: hub covers, explore merchandising, and creator launch assets.

Other roles

Role	Responsibility
Project Manager (Nithila)	Owns the charter, the risk log, open decisions, and the sponsor relationship. Raises things with the Sponsor by exception rather than writing status reports. Single escalation points for blockers. Does not run the technical backlog or assign tasks inside pods.
Sponsor (Adam Kershner)	Approves scope and budget, owns legal due diligence on seeding sources, sends outreach personally, resolves open decisions.

13 Assumptions, Constraints and Dependencies

Assumptions

- Contributors can hold to roughly 20 hours per week across the project.
- Openly licensed material of acceptable quality exists in the tier one topics.
- Enough readers can be recruited for search panel sessions to produce reliable signal.
- Search events and engagement can be linked in Mixpanel or Firebase without major rework.

Constraints

- The initial run is two quarters.
- Work must not disrupt existing creators' publishing or payouts.
- Seeded content must respect license terms, carry attribution, and clear legal review.
- Outreach volume is limited early by sender warm-up, regardless of list size.

Dependencies

- The headline KPI cannot be reported until search success instrumentation is live. Everything else in reporting waits on it.
- Seeding at volume cannot start until the sponsor's legal clearance is in place.

- Outreach depends on domain and Workspace provisioning, with warm-up adding weeks before full volume
- Sourcing priorities stay provisional until panel research runs. Until then, the topic ranking is the working hypothesis.

14 High-Level Risks

- Measuring the wrong thing. Supply counts are easy to move and easy to celebrate. If reporting drifts back to hubs and creators, we can hit every number and still leave readers empty-handed. Mitigated by separating KPIs from operational counts in every report. Owned by the Project Manager.
- Instrumentation slips. Search success rate is the headline KPI and it does not exist yet. If it is not live early, the project runs blind and judges itself on proxies. Owned by the technical group with Analytics.
- Seeding satisfies the metric without serving the mission. Filling topics with public domain titles nobody searches for raises coverage and moves nothing. Mitigated by letting panel findings drive sourcing. Owned by Pod A.
- License and acquisition exposure. Sources mix public domain with permissioned titles, and bulk downloading invites blocking or worse. Mitigated by per-title checks, deliberate pacing, and sponsor-led legal review. Owned by the Sponsor.
- Deliverability collapse. Sending too much too early from new domains burns them and quietly kills outreach. Mitigated by warm-up discipline and volume limits per sender. Owned by Pod B.
- Twenty people, unclear priority. A large team with parallel workstreams can generate activity without progress. Mitigated by self-managing pods with named leads, a single technical owner, and a live KPI dashboard anyone can check. Owned by the pod and workstream leads, escalated by the Project Manager.

15 Risk Register

Risk	Impact	Response
Search success instrumentation delayed	High	Prioritize as the first technical item. Use the ten-query standing set weekly as an interim proxy.
Reporting reverts to supply counts	High	Fixed report template with KPIs first and operational counts in a separate section.
Panel recruitment too slow or unrepresentative	Medium	Start with a small set of sessions, widen sourcing, and supplement with query log analysis once instrumented.
License breach or IP block on a seeding source	High	Per-title verification, no bulk downloading, sponsor-led legal review before volume.
Sender domains burned by early volume	Medium	Five senders, five to ten emails per day each, scale only as reputation builds.
Prospect list too thin to feed outreach	High	Sourcing runs from week one as its own stage with a 200-prospect gate before volume sends. Track list health weekly and use technical automation to enrich rather than building by hand.
Technical backlog becomes a queue with no priority	Medium	Every request states the reader problem and target KPI. Weekly triage owned by Divya, sequenced by KPI impact.
Time zone split slows technical delivery	Low	Schedule handovers around the India time zone member and keep async notes in Linear.
Availability below 20 hours in practice	Medium	Confirm in week one, resize pods and pacing rather than letting scope silently slip.

16 Communication Protocol

Async by default. If something can be a written update, it is not a meeting.

Meetings

For most contributors this is one 30-minute meeting per week, and nothing else.

Meeting	Who	Length	Frequency
Pod standup	Each pod, run by its lead	30 minutes	Weekly (If required)
Technical triage	Technical group, run by Divya	30 - 45 minutes	Weekly (If required)

There is no standing steering meeting. Decisions, risks and blockers go to the Sponsor as they arise rather than waiting for a scheduled slot. The technical triage uses the meeting the technical team already holds, so nothing is added to their calendar. One-to-ones and ad hoc calls happen as needed and are not scheduled ceremony.

Written updates

Update	Owner	Format and timing
Time Log	Each member	The existing company requirement, unchanged. Leads must include a small update about their pods as well as their own tasks for the week, which makes it the pod update as well. Friday end of day.
KPI dashboard	Analytics pair	Live and self-serve refreshed weekly. This is the project's status. Anyone including the Sponsor can look at it at any time without asking for a report.

Reporting by exception

The Project Manager does not write a weekly report. Instead, the dashboard stands as the record, and the Project Manager raises something with the Sponsor only when it needs attention.

The reason the dashboard matters, and Time Logs alone are not enough: Time Logs describe what people did. This project is judged on whether searches succeed, not on how busy the team looked. Without the KPIs visible in one place, the project quietly reverts to being measured by activity, which section 3 exists to prevent.

Panel research findings go straight into a running failed-query log that feeds sourcing priorities. No separate write-up.

If a written update is being duplicated in a meeting, drop the meeting. If a meeting is generating no decisions, drop it and let the leads run async.

17 Open Decisions

Decision needed	What it blocks
Target for search success rate, and the baseline it is measured from.	Whether the 80 percent heuristic is the goal for this run or a longer-term aim.
Definition of a meaningful engagement for the search success metric.	Instrumentation spec and reporting.
How many panel sessions constitute enough signal to redirect sourcing.	When research findings override the topic ranking.
Whether search results are indexed by search engines, given the site blocks crawlers.	How honestly discoverability can be pitched in outreach.
Confirmed weekly availability per contributor against the 20-hour expectation.	Pacing and pod sizing.

18 Change Control and Governance

- Cadence: one weekly pod standup per pod and the technical team's existing weekly triage. No standing steering meeting.
- Escalation: leads raise blockers to the Project Manager as they arise rather than waiting for a meeting. The Project Manager takes anything needing a decision above the team to the Sponsor within 48 hours.
- Reporting: the live dashboard is the project's status, leading with the KPIs in section 3. Operational counts appear separately and are never presented as success.
- Change process: scope or timeline changes above an agreed threshold go to the sponsor before work proceeds.
- This version goes to the full team for review and comment before it is finalized.

19 Approvals

Role	Name	Signature	Date
Sponsor	Adam Kershner		
Technical Product Manager	Divya Emmadi		
Project Manager	Nithila		